

Math 50 Final – Practice

Instructor: Ethan Levien

November 7, 2025

Name: _____

Section: _____

Instructions

- You have 3 hours to complete the exam.
- You may have a one page (single-sided) “cheat sheet” which must be turned in with the exam, but no electronics (including calculators).
- Each problem is worth 5 points.
- Write your solutions in the boxes.
- Don’t cheat.

Exercise 1 (Converting code to math): Consider the following code

```
x1 = np.random.normal(0,1,100)
x2 = 2*x1 + np.random.normal(0,1,100)
y = x1 + x2 + np.random.normal(0,1,100)
b1 = np.cov(y,x1)/np.var(x1)
b2 = np.cov(y,x2)/np.var(x2)
```

What are b1 and b2 (approximately)?

Solution:

Exercise 2 (Joint distribution): Consider the model of a time series X_1, X_2, X_3, \dots :

$$X_{i+1}|X_i \sim \text{Normal}(1 + X_i/2, 1/3)$$

- (a) Write a Python function `generatesim(L)` to generates a simulation of L steps of this time series starting with $X_i = 0$. The function should return a length L numpy array.

- (b) After many steps, the process reaches a steady state where $E[X_{i+1}] = E[X_i]$. What is the distribution of X_i in steady-state?

Exercise 3 (Regression model comparison): (X, Y) data is fit to a single-predictor regression model in statsmodels using OLS, yielding the following output:

	coef	std err	t	P> t	[0.025	0.975]
const	1.0685	0.186	5.756	0.000	0.700	1.437
x1	1.9622	0.177	11.062	0.000	1.610	2.314

A second predictor X_2 is then included in the model, which yields the following output:

	coef	std err	t	P> t	[0.025	0.975]
const	1.0001	0.011	92.003	0.000	0.979	1.022
x1	0.9810	0.012	82.470	0.000	0.957	1.005
x2	2.0122	0.012	168.877	0.000	1.989	2.036

If X_2 and X_1 were fit to a linear regression model with X_2 as the response variable, what would be the regression slope?

Solution:

Exercise 4 (Sample distribution): Consider the model

$$Y_1 \sim \text{Normal}(\beta_1 X_1 + \beta_2 X_2, \sigma_\epsilon^2)$$

After fitting the model, we find $\hat{\beta}_1 = 100$, $\hat{\beta}_2 = -101$, $\hat{\sigma}_\epsilon^2 = 1/4$. The model is then fit to a different data set and it is found that $\hat{\beta}_1 = -100$, $\hat{\beta}_2 = 100.4$.

- (a) Is $\text{cov}(X_1, X_2)$ likely to be positive or negative for the fitted data?

Solution:

- (b) Is it possible that R^2 is very large for this fitted model?

Solution:

Exercise 5 (Bernoulli regression model): Consider the two predictor linear regression model with an interaction term:

$$Y = \beta_1 X_1 + \beta_2 X_2 + J_{1,2} X_1 X_2 + \epsilon$$

The following plot shows data generated from such a model.

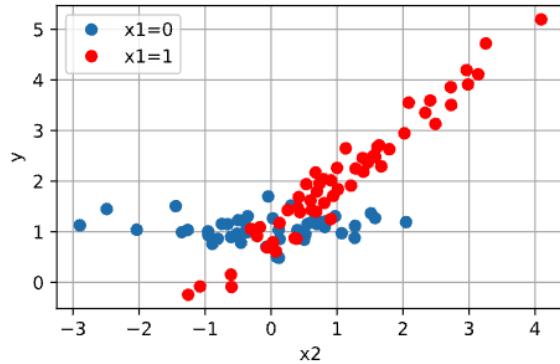


Figure 1:

What (approximately) are the values of β_1 , β_2 and $J_{1,2}$?

Solution:

Exercise 6 (Orthogonality): Consider the features $\phi_1(x) = \sin(2\pi x)$ and $\phi_2(x) = x^2$.

- (a) Are these orthogonal with respect to $X \sim \text{Uniform}(-1, 1)$?

- (b) What is an example of a distribution for X such that ϕ_1 and ϕ_2 are not orthogonal?

Exercise 7 (Missing data): You are given data with predictors X_1, X_2 . You want to fit a linear regression model

$$Y = \beta_0 + \beta_1 X_1 + \beta_2 X_2 + \epsilon$$

but some of the X_2 values have been corrupted and are not reliable. One idea to handle this is called *imputation* and involves generating the missing values X_2 .

- Explain how you could implement imputation assuming you are given a dataframe with rows Y, X_1, X_2, C where $C = 1$ for the rows with corrupted data and $C = 0$ otherwise.
- What issues does your approach to imputation introduce and how might we address these?

Solution: